Taaran Jain

AI / ML Engineer · LLMs, RAG & Agentic Systems

Jaipur, Rajasthan, India | +91 89555 35949 | taaranjain16@gmail.com | taaranjain.com | linkedin.com/in/taaran-jain | github.com/Taaran18

SUMMARY

AI/ML engineer with a year of industry experience building LLM tools for marketing teams, and seven projects that are live online, including a RAG assistant, a churn model and a log anomaly detector. I work mostly in Python, LangGraph and FastAPI, and care about the parts that make AI usable: answers that point to their source, scores that mean something, and systems that keep running when an API fails.

TECHNICAL SKILLS

GenAI & LLMs: LangChain, LangGraph, RAG, Embeddings, Vector Databases (ChromaDB, FAISS), BM25, AI Agents, Prompt Engineering, Structured Outputs, LLM Evaluation, Transformers (Hugging Face)

Machine Learning: Deep Learning, scikit-learn, LightGBM, XGBoost, PyTorch, TensorFlow, Keras, NLP, Anomaly Detection, Probability Calibration

Backend & MLOps: Python, FastAPI, REST APIs, SSE, WebSockets, Docker, PostgreSQL, MongoDB, Supabase, n8n, GCP, AWS, Azure

Frontend: Next.js, React, TypeScript, Tailwind CSS

EXPERIENCE

AI Engineer, e-Marketing.io · Jaipur — Jun 2025 – Present

- Build and ship the agency's AI tools, scoping each problem with the account team, picking the model and deploying it.
- Built an LLM **meeting summarizer** that turns hour-long calls into a brief with decisions and action items, an NLP **keyword analysis engine**, and **lead-qualifying chatbots** that reply to client leads around the clock.
- Built a **WhatsApp task delegation** system (text or voice note) and a **lead management dashboard** in FastAPI, React and PostgreSQL so clients can see where every lead stands.

Data Science Intern, LogiScope Pvt. Ltd. · Bengaluru — Dec 2024 – Jun 2025

- Built anomaly detection for large volumes of system logs with **DBSCAN and a deep learning autoencoder**, flagging problems as they happened and cutting the time to spot an issue by **20%**.
- Wrote the **FastAPI pipeline** around it (log ingestion, preprocessing, inference and alerts) and connected it to the production monitoring setup.

PROJECTS

Nexus AI · *LangGraph, FastAPI, Next.js, SSE* — **Live Demo** · **Case Study**

- Chat assistants rarely show where an answer came from, which makes them hard to trust, and the first version of this one broke overnight when its free LLM was retired.
- Rebuilt it as a six-step **LangGraph** pipeline that chooses between web search, the user's documents (via **embeddings**) or a direct answer, and streams each step as it runs.
- Document answers now cite the exact **file, page and line**. An 18-score decision board comes back in **about 0.4 s**, and cold start fell from up to 1.6 s to **0.7 s**.

DecideX · *scikit-learn, LightGBM, LLMs, FastAPI* — **Live Demo** · **Case Study**

- A churn score on its own doesn't tell a retention team what to do, and with only 26.5% of customers leaving, a model can look accurate while missing most of them.
- Built a tool where an **LLM reads account notes**, a **calibrated model** scores the risk, and a decision step suggests the action and owner, sending unsure cases to a person.
- On 1,409 held-out customers it reached **ROC-AUC 0.844**. Setting the threshold by the cost of a missed churner raised recall to **85.6%**, so missed churners went from 184 to **54**.

Logvera · *scikit-learn, RAG, FastAPI, Next.js* — **Live Demo** · **Case Study**

- Working out what went wrong from a raw server log usually means hours of grep, and a single anomaly model on its own misses a lot.
- Built a log analyser where **five detectors vote** on each time window, and an LLM answers questions using only redacted data and **BM25-retrieved lines** it has to cite.
- On a 34,492-line sample it flagged **3 incidents** and **588 credential-stuffing attempts**. Each upload is encrypted with its own key and deleted after an hour.

EDUCATION

B.Tech, Artificial Intelligence and Data Science, Poornima Institute of Engineering & Technology · CGPA 8.5 — Nov 2021 – Jun 2025

CERTIFICATIONS

MLOps, TensorFlow Developer and Machine Learning Specializations (DeepLearning.AI); WorldQuant Silver; Azure AI Fundamentals

LEADERSHIP

AI/ML Lead, GDSC Poornima (20+ members, 6+ workshops); Event Lead, Poornima Hackathon 2.0 (200+ participants)